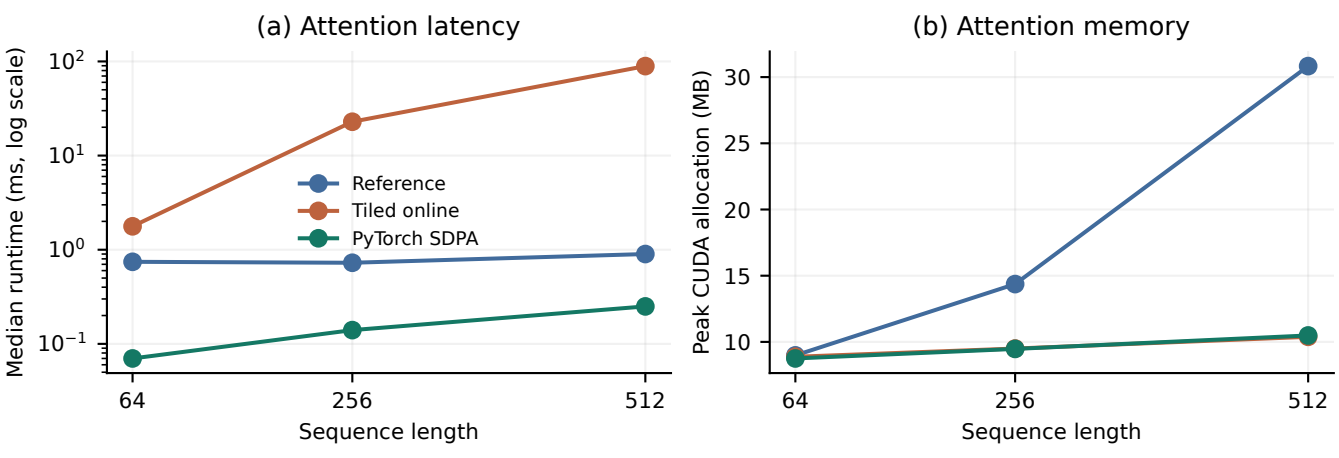
(a) Attention latency
Median runtime (ms, log scale)
10^2
10^1
10^0
10^{-1}
Reference
Tiled online
PyTorch SDPA
64
256
512
Sequence length
(b) Attention memory
Peak CUDA allocation (MB)
30
25
20
15
10
64
256
512
Sequence length