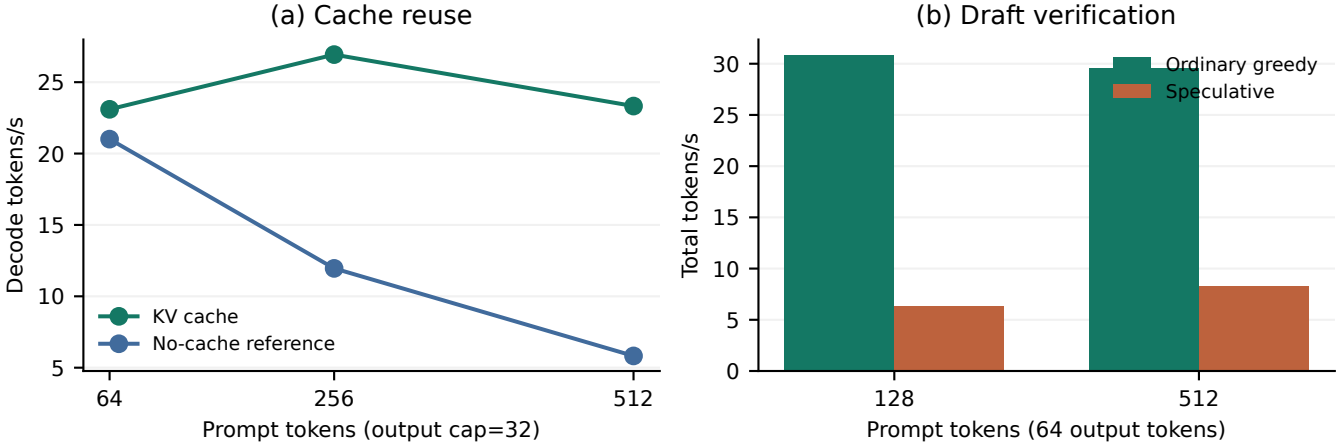
(a) Cache reuse
Decode tokens/s
25
20
15
10
5
KV cache
No-cache reference
64
256
512
Prompt tokens (output cap=32)
(b) Draft verification
Total tokens/s
30
25
20
15
10
5
0
Ordinary greedy
Speculative
128
512
Prompt tokens (64 output tokens)