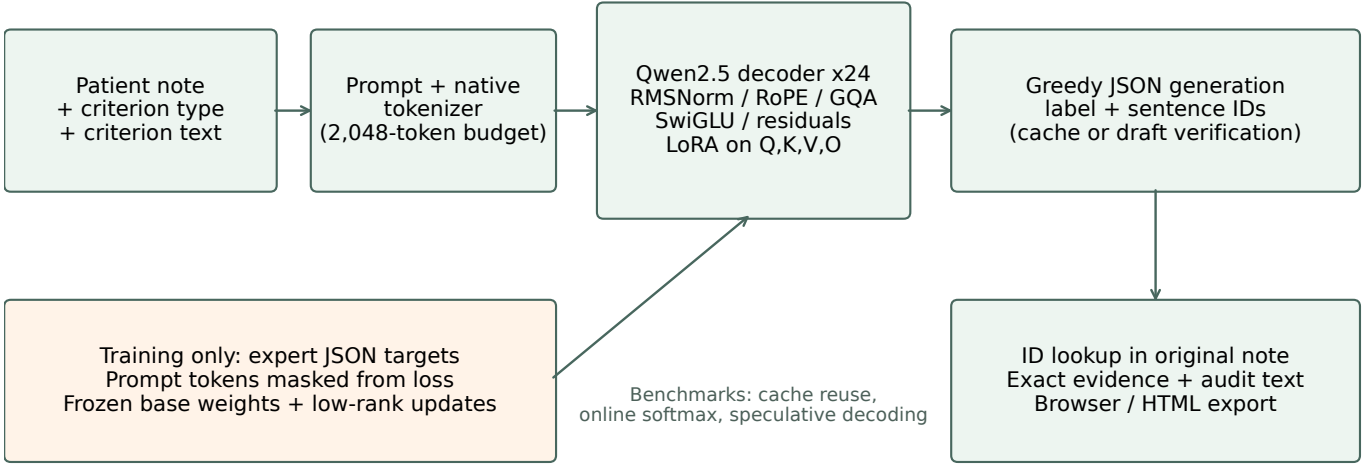
Patient note
+ criterion type
+ criterion text
Prompt + native
tokenizer
(2,048-token budget)
Qwen2.5 decoder x24
RMSNorm / RoPE / GQA
SwiGLU / residuals
LoRA on Q,K,V,O
Greedy JSON generation
label + sentence IDs
(cache or draft verification)
Training only: expert JSON targets
Prompt tokens masked from loss
Frozen base weights + low-rank updates
Benchmarks: cache reuse,
online softmax, speculative decoding
ID lookup in original note
Exact evidence + audit text
Browser / HTML export